

# THE COST OF CONTINUITY

## ROI & PRODUCTION ECONOMICS

Why Generative AI's "Amnesia" is draining studio budgets — and how Mnemosyne fixes it.

**The hidden tax of AI production.** Generating a video or sequence with AI can be cheap. Fixing continuity errors (*style drift, character inconsistency, hallucinated props*) is expensive. Studios are often misled by low API costs while paying a large, ongoing invoice for **human-in-the-loop rework**.

## COST COMPARISON: 100-SHOT SEQUENCE (Approx. 60 seconds)

Assumption: Standard commercial rate for a VFX/AI Artist is \$100/hour.

Metric	BASELINE: Standard GenAI Pipeline (Stateless)	MNEMOSYNE PIPELINE (ICF + Verifier)
Initial generation API cost	~\$10.00	~\$10.00
Verifier compute	—	+\$2.50 (local/heuristic; shadow gating)
Continuity hallucination rate	60% (60 shots require fixing)	Reduced; edge-cases caught pre-render
Manual rework & inpainting	15 artist-hours (lighting, faces, outfits)	2 artist-hours (only extreme edge-cases)
Total human labor cost	\$1,500.00	\$200.00
Total time to delivery	3–4 days	1 day
<b>BOTTOM LINE</b>	<b>\$1,510.00 (plus schedule slip risk)</b>	<b>\$212.50 (on-time delivery)</b>

## THE BUSINESS VALUE (ROI)

- 85% reduction in rework costs.** Mnemosyne shifts continuity enforcement from expensive post-production artist hours to inexpensive mid-production verification compute.
- Scalable IP governance.** Mnemosyne integrates into existing pipelines (Nuke, Premiere, ShotGrid) as a background orchestration layer — without forcing a new UI or per-seat licensing model.
- Model-agnostic and future-proof.** Studios don't need to abandon Midjourney or Sora. As foundational models evolve, Mnemosyne remains the persistent "security and continuity rail" underneath them.

*Note:* All figures are illustrative and intended for investor data-room communication. Actual rates and savings vary by studio, content type, and tooling stack.