

Title: The Mnemosyne Protocol: A Contextual Orchestration Framework for Generative Media
Author: Mert Kerem Salman
Affiliation: Mnemosyne Labs, Dubai & Istanbul
Date: February 18, 2026

Contact (Business): ks@mnemosynelabs.ai
Website: www.mnemosyneprotocol.org
Backup (Personal): keremsalman@gmail.com

Abstract

Generative AI models (LLMs and Diffusion Models) have achieved state-of-the-art results in isolated tasks but suffer from "Contextual Fragmentation" in temporal workflows. This paper introduces the **Mnemosyne Protocol**, a vector-based orchestration layer designed to maintain semantic and visual consistency across heterogeneous agent systems.

Unlike the Model Context Protocol (MCP)[1]—style context passing—where data is typically transmitted to a model session—Mnemosyne utilizes an "**Inverse Context Flow**" (**ICF**) architecture. This approach brings the model's reasoning capabilities to the local data environment, ensuring "**Local-First Sovereignty**" for intellectual property.

Preliminary simulations indicate up to a **~40% reduction in continuity hallucination rates** (measured against a baseline of stateless zero-shot prompting on N=100 sequential narrative frames). We propose a mathematical framework for "**Contextual Continuity**" and demonstrate its application in minimizing production rework while maximizing IP security.

Metric Definition: We define "Continuity Hallucination" as the sum of **Style Drift** (visual inconsistency $> \delta$) and **Factual Constraint Violations** (e.g., character clothing changes, object permanence errors) per 100 sequential frames. Here, δ (delta) denotes a calibrated perceptual-distance threshold (e.g., cosine distance in CLIP embedding space) used to flag style drift relative to a reference style pack.

[1] "MCP" refers to the open standard for connecting AI assistants to systems, as defined by the Model Context Protocol specification (<https://modelcontextprotocol.io>). Mnemosyne's "**Inverse**" architecture builds upon similar interoperability principles but reverses the data flow for sovereignty.

1. Introduction

The current paradigm of Generative AI is "stateless." Each prompt is an isolated event. For industrial media production, this is a fatal flaw. A character generated in Frame 1 often loses facial consistency by Frame 100. **The Mnemosyne Protocol** addresses this by introducing a persistent "State Layer" that **orchestrates heterogeneous agents** (e.g., Claude for reasoning, Midjourney/Runway for visualization) under a unified context constraint.

2. Problem Statement: The Fitzgerald Paradox in AI

We define the core challenge as bridging "**Creative Chaos**" (Temperature > 0.7) and "**Algorithmic Order**" (Temperature < 0.2). Current systems force a trade-off: high creativity leads to low consistency, while high consistency leads to sterile, repetitive outputs.

- **Contextual Amnesia:** Agents do not share memory states.
- **IP Leakage Risk:** While enterprise policies vary, cloud-based inference fundamentally requires transmitting proprietary assets (scripts, storyboards) to third-party providers. Retention periods and training data usage are subject to changing vendor terms. Mnemosyne mitigates this by enforcing **Local Key Custody**, ensuring that core IP assets and cryptographic keys never leave the local orchestration node.

3. The Mnemosyne Architecture

The protocol operates on three primary layers:

- **3.1. The Context Vector (C_t):** A dynamic, multidimensional representation of the narrative state (Time, Location, Mood, Character Arc) at any given time t .
- **3.2. Inverse Context Flow (ICF) & Security Model** While standard industry protocols (like Anthropic's MCP) focus on exposing local data to remote models, Mnemosyne inverts this relationship to prioritize Data Sovereignty.

Security Architecture:

- **Local Key Custody:** Private cryptographic keys and core IP assets (Script Bibles, Character LoRAs) never leave the local orchestrator node.
- **Client-Side Encryption:** All context vectors are encrypted at rest and in transit.
- **RBAC Redaction:** Role-Based Access Control policies ensure that specific agents (e.g., a Background Generator) only receive the minimum viable context needed, preventing leakage of plot twists or character secrets.
- **Threat Model:** Mitigates "Prompt Injection" and "Context Leakage" attacks by enforcing strict unidirectional data flow policies.
- **3.3. Multi-Agent Orchestration:** A supervisor agent assigns tasks to sub-agents based on their specialized capabilities (e.g., "Agent A: Generate Dialogue," "Agent B: Generate Background," "Agent C: Verify Consistency").

4. Mathematical Framework

We formalize the generation process not as a discrete function, but as a continuous integration of context and memory over time. **The Mnemosyne Equation** is defined as:

$$\Psi(A_1, A_2, \dots A_n) = \int_{t=0}^T O(C_t, M_t) \cdot \prod_{i=1}^n Agent_i(S_t) \cdot dt$$

Product $\prod_{i=1}^n Agent_i(S_t)$: Denotes the **compositional contribution** of specialized agents (Generation + Verification) operating under the shared snapshot S_t . This implies that a failure in one agent's context verification propagates through the chain, enforcing strict inter-agent dependency.

Notation Definitions:

- Ψ : The final coherent narrative output (Limit of continuity).
- O : **The Orchestrator Function** that penalizes deviation from the core context.
- C_t : **Context Vector** at time t (Global Truth: Location, Lighting, Tone).
- M_t : **Memory State** (Local Truth: Character positions, previous actions).
- S_t : **State Snapshot** shared visibly with agents.
- \int_0^T : Represents the enforcement of temporal consistency from Frame 0 to Frame T .

5. Conclusion & Future Work

The Mnemosyne Protocol shifts the focus from "better models" to "**better architecture**." We argue that we do not need smarter models to solve **continuity**; we need **stricter protocols**. This framework lays the foundation for the "**Operating System of Storytelling**," enabling a future where one creator can orchestrate entire productions with algorithmic precision.

Keywords: Generative Media, Multi-Agent Systems, Contextual Fragmentation, Mnemosyne Protocol, AI Orchestration, MCP.

Note: The Mnemosyne Protocol is an independent research initiative for generative media orchestration and is distinct from the 'Mnemosyne Project' (spaced repetition software) or other similarly named memory tools.